Responsible AI Use & Doctoral Academic Integrity

Submitted for Course Assignment (50 points)

Due: Friday, October 2, 2026

AI Disclosure

AI tool(s) used: Claude (Sonnet 4.5), accessed through the Cursor platform.

Complete prompts/instructions provided: Across this term's coursework, I used this tool to: (1) help draft an original discussion post on the tension between privacy and convenience in digital technology, including peer replies; (2) check specific phrases from that post against publicly available web sources to confirm I had not inadvertently duplicated existing text; (3) help me understand and respond to this course's AI Use Policy across several drafts, including the present assignment, which I provided in full along with its four-part instructions; (4) help identify relevant data points and construct a full dimensional data model (fact and dimension tables) for a university data warehouse scenario; (5) answer clarifying questions (e.g., defining the term "LMS") and verify word counts of drafted text; (6) help draft and revise replies to classmates' discussion posts on data warehousing's impact on institutional decision-making; and (7) for this specific assignment, locate and verify three scholarly sources on AI ethics and academic integrity, draft the policy analysis in Part II, generate the AI appendix response and critical evaluation in Part III, and draft the personal standard in Part IV.

What I used the AI tool to accomplish: Primarily to produce first-draft content across multiple assignments, discussion posts, data model designs, peer replies, and this reflection itself, based on prompts I provided, and secondarily to verify factual claims, word counts, and source accuracy.

Which portions of the AI output influenced my final work: Substantially all of it. Across this term, I largely requested complete drafts rather than producing my own first attempt and seeking AI feedback on it afterward. This is addressed directly and critically in Part II and Part IV of this submission, where I argue that this pattern represents the kind of intellectual substitution this policy is designed to flag, and where I commit to a different sequencing of my own work versus AI assistance going forward.

Part II: AI Policy Analysis

The course AI Use Policy is built around a distinction that is easy to state and hard to apply consistently: AI may support thinking, but it may not replace it. From the perspective of a doctoral scholar, this distinction is not a bureaucratic formality. It is the mechanism by which the credential itself retains meaning.

Disclosure matters because it is the only practical substitute for a trust relationship that no longer exists by default. In traditional scholarship, a reader could reasonably assume that prose reflects the author's own reasoning unless told otherwise. Generative AI breaks that assumption silently; text produced with substantial AI assistance can be indistinguishable, on the page, from text produced through independent analysis. Eaton (2023) argues that this is precisely why detection-based approaches to academic integrity are approaching obsolescence: as AI and related technologies become normal, invisible parts of how people think and write, policing originality after the fact becomes increasingly unworkable. If detection cannot reliably distinguish AI-assisted work from independent work, the burden necessarily shifts to disclosure. Disclosure is not a punitive mechanism; it is what allows a reader, an advisor, or a dissertation committee to evaluate what kind of intellectual labor actually produced a given piece of work, which is the only way the underlying claim of doctoral scholarship, that the degree certifies independent capability, remains verifiable at all.

The distinction between AI supporting thinking and AI replacing thinking turns on sequencing and authorship of ideas rather than on whether a tool was used at all. Using AI to clarify a term, check a citation format, or identify a weakness in an argument I have already constructed leaves the locus of intellectual production with me; I am the one who generated the claim, and the tool is reacting to it. Using AI to generate the claim itself, the literature synthesis, or the analytical framework, and then presenting that output as my own interpretation, relocates the locus of production to the tool. Chan and Hu (2023), in their survey of university students' attitudes toward generative AI, found that students themselves recognized this tension: while most valued AI for brainstorming and writing support, many also expressed concern about overreliance eroding their own development as independent thinkers. That concern is amplified, not reduced, at the doctoral level, where the entire purpose of the degree is to certify that independent capability.

Excessive reliance on generative AI in doctoral scholarship carries at least two distinct ethical risks. The first is epistemic: generative AI models produce fluent, confident text regardless of whether the underlying claims are accurate, and they are known to fabricate citations and misrepresent sources. A doctoral scholar who relies on AI output without independently verifying every factual and bibliographic claim risks introducing fabricated evidence into the scholarly record under their own name. The second risk is developmental: doctoral training is not only about producing a dissertation but about developing the capacity for independent

critical analysis that will be exercised long after the degree is conferred. Eaton's (2023) broader argument about equity in a "postplagiarism" era is relevant here as well, since unequal access to AI tools, and unequal AI literacy, could translate into unequal opportunities to either benefit from or be harmed by these same dynamics across different scholars and institutions.

Finally, the policy's requirement that responsibility remains with the researcher is not merely a liability-shifting clause; it reflects the only coherent allocation of accountability available once a non-human tool is involved in producing scholarly output. Floridi and Cowls (2019) identify explicability, encompassing both intelligibility and accountability, as a necessary principle precisely because AI systems cannot themselves be held responsible for the accuracy or ethical standing of their output. An AI tool cannot be cited as an authority, cannot defend an interpretation under scrutiny, and cannot be held to a standard of scholarly integrity, because it has no professional standing to lose. Responsibility therefore cannot be distributed to the tool; it remains entirely with the scholar who selected, prompted, evaluated, and ultimately submitted the work under their own name. This is why the policy's framing of AI output as "a starting point, not the finished intellectual product" is the correct standard: the researcher is accountable not for the fact that a tool was consulted, but for every claim, citation, and conclusion that ultimately appears in work bearing their name.

References

Chan, C. K. Y., & Hu, W. (2023). Students' voices on generative AI: Perceptions, benefits, and challenges in higher education. International Journal of Educational Technology in Higher Education, 20, 43. https://doi.org/10.1186/s41239-023-00411-8

Eaton, S. E. (2023). Postplagiarism: Transdisciplinary ethics and integrity in the age of artificial intelligence and neurotechnology. International Journal for Educational Integrity, 19(1), 1-10. https://doi.org/10.1007/s40979-023-00144-1

Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. Harvard Data Science Review, 1(1).

Part III: Evaluate the AI

Appendix - AI Response to the Prompt:

"What are the most important ethical principles a doctoral researcher should follow when using generative AI in academic research?"

1. Transparency and Disclosure: Always disclose when and how AI tools were used in your research process, including which tasks AI assisted with.

2. Accountability: You remain responsible for all content in your work, including anything generated or assisted by AI.

3. Accuracy and Verification: AI tools can produce incorrect information or "hallucinate" facts and citations. Researchers should independently verify any factual claims before including them in academic work.

4. Avoiding Plagiarism and Ensuring Originality: AI-generated text should not be presented as entirely original work without disclosure.

5. Respecting Privacy and Confidentiality: Researchers should be cautious about what information is shared with third-party AI platforms.

6. Maintaining Academic Rigor: AI should enhance, not replace, critical thinking and original argument development.

7. Fairness and Bias Awareness: AI models can reflect biases present in their training data, especially in analysis involving social, cultural, or demographic factors.

By following these principles, doctoral researchers can use generative AI as a valuable tool while maintaining the integrity, rigor, and ethical standards expected of academic scholarship.

Critical Evaluation

The AI response correctly identifies several principles that align with established frameworks in the scholarly literature. Its emphasis on transparency and disclosure maps onto the same accountability logic Floridi and Cowls (2019) describe as "explicability," and its warning about AI hallucinating facts and citations reflects a well-documented and genuine risk rather than an invented concern. The response is also correct that responsibility for errors remains with the researcher rather than the tool. As a baseline inventory of relevant concerns, the response is reasonably accurate and not misleading in any specific claim it makes.

What it oversimplifies, however, is more significant than what it gets right. The response presents seven principles as a flat, equally weighted checklist, with no indication of which tensions are hardest to resolve in practice. It never engages with the actual difficulty at the center of this assignment: distinguishing AI as support for thinking from AI as a substitute for thinking is not a bright line but a judgment call that depends on sequencing and intent, which a

numbered list cannot capture. It also treats "accuracy and verification" as a simple technical check, when in practice verifying an AI-generated claim often requires the same disciplinary expertise a doctoral student is still developing, precisely what makes over-reliance risky.

Several important perspectives are missing entirely. The response says nothing about equity, despite the fact that access to capable AI tools and institutional AI literacy varies considerably across institutions, a concern Eaton (2023) raises directly in her discussion of a "postplagiarism" era. It also omits the developmental purpose of doctoral training: the response treats AI ethics as risk management rather than acknowledging that the dissertation process exists to build the researcher's own independent capability, which over-reliance undermines even absent any factual error. It does not address discipline-specific variation either, though what counts as appropriate AI assistance differs meaningfully between quantitative and interpretive research traditions.

The response provided no citations, sources, or references of any kind, despite making empirical claims, for example, that AI models "can reflect biases present in their training data." This is itself instructive: a generic AI response to an ethics question will often assert accurate-sounding claims without supporting evidence, which is exactly the verification burden the policy places on the researcher. Every claim in the appendix would need independent sourcing before being relied upon in scholarly writing.

To improve the response, I would restructure it around the actual tension doctoral scholars face, assistance versus substitution, and explicitly integrate the sources used in Part II: Floridi and Cowls' explicability framework to ground accountability, Eaton's postplagiarism argument to address equity and the limits of detection-based enforcement, and Chan and Hu's empirical finding that students themselves worry about overreliance eroding their own development. That synthesis would turn a generic ethics checklist into an argument specific to the stakes of doctoral scholarship.

Part IV: Personal AI Integrity Standard

Having spent this assignment analyzing the policy and critiquing an AI-generated response, I want to locate the line between assistance and substitution somewhere more precise than "it depends." The line, for me, is this: AI assistance preserves my authorship of the claim, while intellectual substitution transfers authorship of the claim to the tool, regardless of how much editing I do afterward. If I have already formed a position, an interpretation, or an argument, and I use AI to test it, sharpen its expression, or check it against a counterargument, I remain the author. If I ask AI to form the position for me and then revise its wording, the tool has authored the thinking even though my name appears on the page. Polished prose can disguise this distinction entirely, which is exactly why the policy's disclosure requirement matters: it forces me to be honest about which side of that line a given piece of work actually falls on, independent of how finished it looks.

Based on this standard, I am committing to three specific practices for the remainder of my doctoral studies. First, I will produce an independent first attempt at any analysis, however rough, before introducing AI into the process at all; this guarantees that the core claim originates with me and that AI can only ever refine or challenge it, never generate it in my place. Second, I will treat every factual claim, citation, or source that AI surfaces as unverified until I have checked it against the original material myself, since the epistemic risk of fabricated or misrepresented evidence does not diminish just because the prose sounds confident. Third, I will disclose AI use proportionally to its actual role in the final work, meaning a detailed account of what the tool did and which parts of its output survived into my submission, rather than a vague or minimal acknowledgment that technically satisfies a requirement without reflecting what actually happened.

I recognize that these commitments will be tested most by convenience rather than by any single dramatic temptation to cheat. The real risk this term has shown me is incremental: a fast, competent tool makes it easy to lean on it slightly more with each assignment until the habit of producing my own first draft quietly disappears. Holding the line I have described here means noticing that drift before it becomes normal, and treating the extra friction of thinking first as a cost worth paying, because it is the actual work this degree exists to certify that I can do.